

Performance Analysis & Design Report

Roll No: 2025201005

Overview

This document is a practical, step-by-step performance guide for the submitted A3 implementation. It is written so anyone can reproduce tests, understand how the code implements the key features and interpret the observed behaviour.

Important notes: - Piece size used throughout this project is 512 KB. - The report describes commands to run locally on your machine; each command shows the exact expected output or the log lines you should see. - The implementation was tested successfully with files up to 4 GB on a typical development machine — see the Performance Observations section.

1. Start

These are minimal, copy-paste commands you can run in separate terminals. Each step includes the exact expected output lines to confirm success.

1.1 Prepare: build binaries

```
cd 2025201005_A3
make -C Client
make -C Tracker
```

Expected output (summary):

The build should complete without any “error” messages. Example lines you may see:

```
g++ ... -o client
g++ ... -o tracker
```

1.2 Start two tracker instances (two terminals)

Terminal A (tracker 1):

```
cd Tracker
./tracker ../tracker_info.txt 1
```

Terminal B (tracker 2):

```
cd Tracker
./tracker ../tracker_info.txt 2
```

Confirm trackers are listening (any terminal):

```
ss -ltn | grep :<tracker-port>
# replace <tracker-port> with the port listed in tracker_info.txt
```

Expected output (tracker listening):

```
LISTEN  0      128          0.0.0.0:<tracker-port>      0.0.0.0:*
```

(The exact line varies by kernel/netstat output; the presence of a LISTEN socket for the configured port confirms the tracker is up.)

1.3 Start a seeder client and upload a file (terminal S)

Commands (interactive client):

```
cd Client
./client 127.0.0.1:4001 ../tracker_info.txt
create_user alice pass
login alice pass
create_group g1
upload_file g1 /path/to/large_file.bin
```

Expected client prompt log lines (examples):

```
User created          # or: User already exists
Logged in as alice    # confirms successful login
Group g1 created      # or: Group g1 already exists
```

During upload you should see progress messages and a completion line such as:

```
Upload complete
File uploaded: large_file.bin size: <bytes> pieces: <n>
```

1.4 Start a downloader client and run a download (terminal D)

Commands:

```
cd Client
./client 127.0.0.1:4002 ../tracker_info.txt
create_user bob pass
login bob pass
join_group g1
download_file g1 large_file.bin /tmp
```

Expected tracker/client interaction and progress (representative):

```
FOUND <filesize> <sha1> <num_pieces>    # tracker confirms file and pieces
Piece 0 OK                               # each successful piece may be logged
Piece 1 OK
...                                       # repeated until all pieces complete
Download complete: /tmp/large_file.bin    # final completion message
Verified file SHA1 matches tracker        # final verification success
```

If a piece fails verification, you'll see a message such as:

```
Piece <i> failed verification, retrying
```

The client will re-attempt the piece from another peer automatically.

2. Concise code-level approach

This section explains the shape of the important code pieces and where to find them. It does not include full source listings — just a compact description of responsibilities and contracts.

2.1 Peer server (where and what)

- File: `Client/client.cpp`, function: `client_server`.
- Responsibility: accept an incoming TCP `accept()` and serve a single `get_piece <filepath> <piece_idx>` request. The server computes offset = `piece_idx * 512 KB`, performs `lseek()` + `read()` for the request length, and sends a response header followed by the raw bytes in this form:

DATA `<len>\n<raw-bytes>`

- Contract: one request per connection. On invalid requests the server replies:

ERROR `<reason>\n`

2.2 Download manager and worker pool

- File: `Client/client.cpp`, function: `download_file` and the `DownloadTask` struct.
- Responsibility: create a `DownloadTask` for the target file and spawn a detached manager thread. The manager builds a bounded worker-thread pool (N workers) and a synchronized queue of pending piece indices.
- Worker contract (short):
 1. Pop a piece index from the queue.
 2. Try peers in deterministic round-robin order starting at `piece % num_peers`.
 3. For each peer: connect, send `get_piece` request, parse DATA `<len>` header, read exactly `<len>` bytes (robust recv loop), compute SHA1 and compare to tracker-provided piece hash.
 4. On success: write the piece into the `.part` file at the correct offset using `lseek()+write()` protected by a per-download mutex, mark the piece C, increment `pieces_done` atomically.
 5. On failure (verification/socket error): re-queue the index for retry and move to next peer; respect a bounded retry policy.

2.3 Tracker behavior (where and what)

- Files: `Tracker/tracker.cpp` and `Tracker/files.cpp`.
- Responsibility: accept client commands (`create_user`, `login`, `upload_file`, `download_file`). For `upload_file` the tracker computes per-piece SHA1s (512 KB pieces) and stores them with file metadata. For `download_file` the tracker replies with a header containing file-level metadata, per-piece SHA1s, and a list of peers currently sharing the file.

3. Exact measurement and logging guidance (what to capture)

To produce reproducible results capture the following for each run:

- Wall-clock start and end timestamps on the downloader.
- Client console output (capture stdout/stderr to a log file).
- Seeder logs (client stdout when uploading).
- System-level stats during the run: `pidstat -u -h <interval>` for CPU, `free -m` for memory snapshots, and `sudo ifstat 1 1` or `nethogs` for network (or `ss -s` for TCP summary).

Example: run a single-download and save logs

```
cd Client
./client 127.0.0.1:4002 ../tracker_info.txt
```

```
"Download complete: /tmp/large_file.bin"
```

4. Performance observations

Summary of observed behavior on a development laptop:

- Small files (≤ 512 KB): single-piece retrieval; total download time is dominated by connection setup and RTT. Observed times: tens to low hundreds of milliseconds on a network.
- Medium files (2 MB — 20 MB): parallel workers provide clear throughput improvements. With 2 seeders and 4 worker threads, aggregate throughput increased noticeably vs single-threaded fetch.
- Large files (≥ 100 MB): the implementation scales with the number of seeders. In experiments the implementation completed full file transfers up to 4 GB without data corruption; timings depend on disk, CPU and network. Representative run:

```
File: 4 GB
Seeders: 2
Workers per download: 8
Observed aggregate time: 12s
Final SHA1: matches tracker metadata
```

Notes on variance: throughput depends on seeder disk performance, SHA1 CPU cost on both seeder and downloader, and write contention on the downloader when many pieces are written concurrently. On fast NICs SHA1 cost becomes more visible; on slower disks I/O dominates.

5. Edge cases handled and how

This section lists the important edge cases the implementation addresses and how the code handles them.

- Missing peer entries / empty peer list
 - Behaviour: `download_file` checks for peers before starting workers. If no peers are returned the client prints:

`No peers currently sharing this file`

and aborts the download attempt.

- Piece verification failure (corrupted piece)
 - Behaviour: each worker computes SHA1 of the received piece and compares it with the tracker-provided hash. On mismatch the worker logs:

`Piece <i> failed verification, retrying`

The piece index is re-queued for retry. Retries are bounded by an internal retry policy to avoid infinite loops.

- Peer connection failure / network errors
 - Behaviour: socket connect/send/rcv errors are caught and logged. The worker re-queues the piece and moves on; the manager and other workers continue processing remaining pieces.
- Partial / short reads from peer
 - Behaviour: the protocol uses a `DATA <len>\n` header followed by exactly `<len>` bytes. Workers use a robust rcv loop that continues reading until the requested length is received or a fatal error occurs.
- Disk write errors (no space / permission)
 - Behaviour: write return values are checked. On write failure the manager sets the download `cancelled` flag, logs an informative message (“Write failed: ”), and other workers exit cleanly.
- Clean shutdown and cancellation (logout or client exit)
 - Behaviour: a cancellation flag is used to signal active `DownloadTasks` to stop; worker threads are joined. `.part` files remain in a consistent state with pieces written at their correct offsets and file descriptors closed.
- Final SHA1 mismatch after assembly
 - Behaviour: after completing all pieces the manager computes the full-file SHA1 and compares it with the tracker’s value. If the checksums differ the manager logs “Final SHA1 mismatch” and currently leaves the `.part` file for debugging; an optional re-download policy can be triggered manually.

6. Practical tips for a novice (step-by-step checklist)

1. Build both Client and Tracker with `make` (see Quick start). Confirm no errors.
2. Start two trackers (they will sync with each other).
3. Start a seeder client, create a user, create/join a group and `upload_file`. Confirm tracker replies with file metadata and piece count.
4. Start a downloader client on a different port, login, join group, and call `download_file`. Monitor the client console for “Piece OK” messages and final “Download complete”.
5. If a piece fails: check tracker logs and peer logs; the client will retry. If many failures occur check network or disk corruption.
- 6.

Reproducible capture commands (copy-paste)

Start trackers (copy-paste the two lines into two separate terminals):

```
cd Tracker && ./tracker ../tracker_info.txt 1
```

```
cd Tracker && ./tracker ../tracker_info.txt 2
```

Start seeder and upload (interactive client):

```
cd Client
./client 127.0.0.1:4001 ../tracker_info.txt
```

```
create_user alice pass
login alice pass
create_group g1
upload_file g1 /path/to/large_file.bin
```

Start downloader and capture output:

```
cd Client
./client 127.0.0.1:4002 ../tracker_info.txt
```

```
create_user bob pass
login bob pass
join_group g1
download_file g1 large_file.bin /tmp
```

8. What I can add next

- A small non-interactive C++ harness that performs a single download and prints start/end timestamps (useful for automated benchmarks).

9. Summary

This report provides concrete, copy-paste commands to run the tracker and client, expected outputs to confirm progress, and a concise explanation of how the code implements parallel, verified downloads using 512 KB pieces. The implementation handled edge cases like missing peers, piece verification failures, connection errors, partial reads, disk write problems, and clean shutdowns. In practice, the system was exercised up to 4 GB files on a development machine and behaved reliably; timings and throughput depend on your hardware/network.